


Google Play Store Apps Report



Data Analyst: Chaimaa Driouech

GitHub:

Date: 01/04/2020

Summary

OVERVIEW	3
Acknowledgements.....	3
Inspiration	3
Tasks.....	3
PYTHON AND ALL IMPORTED LIBRARIES.....	4
EXPLORATORY DATA ANALYSIS	5
DATA COLLECTION AND PRE-PROCESSING.....	10
FEATURE ENGINEERING	15
Transforming the Predictive Target (Y) using Label Uncoder	15
Method 1: Recursive Feature Elimination	15
Method 2: Feature Importance.....	16
DATA VISUALIZATION	18
App with large number of reviews.....	18
App with the largest size	18
App with the largest num of installs.....	18
Paid vs Free	19
App which hasn't been updated	19
Most popular category	20
DATA SCIENCE PROCESS: BEST PRACTICES	22
BADIR: Project Process Framework	22
PREDICTIVE MODELING AND EVALUATION (THE WHOLE PROCESS).....	24
Random Forest Regression.....	26
Support Vector Regression.....	26
Linear Regression	27
Predictions.....	27
MODEL SELECTION	28
CROSS VALIDATION	29

OVERVIEW

While many public datasets (on Kaggle and the like) provide Apple App Store data, there are not many counterpart datasets available for Google Play Store apps anywhere on the web. On digging deeper, I found out that iTunes App Store page deploys a nicely indexed appendix-like structure to allow for simple and easy web scraping. On the other hand, Google Play Store uses sophisticated modern-day techniques (like dynamic page load) using JQuery making scraping more challenging.

Acknowledgements

This information is scraped from the Google Play Store. This app information would not be available without it.

Inspiration

The Play Store apps data has enormous potential to drive app-making businesses to success. Actionable insights can be drawn for developers to work on and capture the Android market!

Tasks

- GitHub
- Python and all libraries needed to solve the problem
- Exploratory Data Analysis
- Data collection, pre-processing and feature engineering
- Data Visualization
- Data science process: Best practices
- Predictive Modeling and Evaluation (the whole process)
- Model selection
- Cross validation

PYTHON AND ALL IMPORTED LIBRARIES

```
#imports
#numpy,pandas,scipy, math, matplotlib
import numpy as np
import pandas as pd
import scipy
from math import sqrt
import seaborn as sns
import matplotlib.pyplot as plt

#estimators Regression
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn import linear_model

#estimators Classification
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVR
from sklearn import svm
from sklearn.neural_network import MLPClassifier
from sklearn.preprocessing import StandardScaler, LabelEncoder

#model metrics Regression
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.model_selection import cross_val_score
#model metrics Classification
from sklearn.metrics import confusion_matrix, classification_report

#cross validation
from sklearn.model_selection import train_test_split

import warnings
warnings.filterwarnings('ignore')
pd.options.display.max_columns = None

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import StratifiedKFold
from sklearn.feature_selection import RFECV

%matplotlib inline
```

EXPLORATORY DATA ANALYSIS

df_apps.dtypes.index

```
Index (['App', 'Category', 'Rating', 'Reviews', 'Size', 'Installs', 'Type',  
        'Price', 'Content Rating', 'Genres', 'Last Updated', 'Current Ver',  
        'Android Ver'],  
      dtype='object')
```

df_apps.info()


```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10841 entries, 0 to 10840  
Data columns (total 13 columns):  
App                10841 non-null object  
Category           10841 non-null object  
Rating            9367 non-null float64  
Reviews           10841 non-null object  
Size              10841 non-null object  
Installs          10841 non-null object  
Type              10840 non-null object  
Price             10841 non-null object  
Content Rating    10840 non-null object  
Genres            10841 non-null object  
Last Updated      10841 non-null object  
Current Ver       10833 non-null object  
Android Ver       10838 non-null object  
dtypes: float64(1), object(12)  
memory usage: 1.1+ MB
```

df_apps.describe()


	Rating
count	9367.000000
mean	4.193338
std	0.537431
min	1.000000
25%	4.000000
50%	4.300000
75%	4.500000
max	19.000000

Content Rating

Out of the 10840 apps stored in our Dataset, 80 % of the applications are targeting all age groups from children, mature 21+ to adult as shown in the graph below.




Type vs Rating




Rating

Overall user rating of the app (as when scraped) shows in the following figure over 1000 apps were rated 4.4 as highest rating.




Genres


An application can belong to multiple genres (apart from its main category). For eg, a musical family game will belong to



Category vs Rating



Content Rating vs Rating



DATA COLLECTION AND PRE-PROCESSING

Data Collection

This information was scraped from the **Google Play Store**. Using the **Web Scraping** technique (also termed *Screen Scraping*, *Web Data Extraction*, *Web Harvesting* etc.) it is employed to extract large amounts of data from websites whereby the data is extracted and saved to a local file in your computer or to a database in table (spreadsheet) format.

Data displayed by most websites can only be viewed using a web browser. They do not offer the functionality to save a copy of this data for personal use. The only option then is to manually copy and paste the data - a very tedious job which can take many hours or sometimes days to complete. Web Scraping is the technique of automating this process, so that instead of manually copying the data from websites, the Web Scraping software will perform the same task within a fraction of the time.

There are 13 features in our Data (as when scraped):

1. **App:** Application name
2. **Category:** Category the app belongs to
3. **Rating:** Overall user rating of the app
4. **Reviews:** Number of user reviews for the app
5. **Size:** Size of the app
6. **Installs:** Number of user downloads/installs for the app
7. **Type:** Paid or Free
8. **Price:** Price of the app
9. **Content Rating:** Age group the app is targeted at - Children / Mature 21+ / Adult
10. **Genres:** An app can belong to multiple genres (apart from its main category). For eg, a musical family game will belong to Music, Game, Family genres.
11. **Last Updated:** Date when the app was last updated on Play Store
12. **Current Ver:** Current version of the app available on Play Store
13. **Android Ver:** Min required Android version

Data Pre-processing

#Detection of Missing Values

```
df_apps.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10841 entries, 0 to 10840
Data columns (total 13 columns):
App                10841 non-null object
Category          10841 non-null object
Rating           9367 non-null float64
Reviews           10841 non-null object
Size              10841 non-null object
Installs          10841 non-null object
Type            10840 non-null object
Price             10841 non-null object
Content Rating  10840 non-null object
Genres            10841 non-null object
Last Updated      10841 non-null object
Current Ver     10833 non-null object
Android Ver    10838 non-null object
dtypes: float64(1), object(12)
memory usage: 1.1+ MB
```

```
df_apps.isnull().sum()
```

```
App                0
Category           0
Rating           1474
Reviews           0
Size              0
Installs          0
Type            1
Price             0
Content Rating  1
Genres            0
Last Updated      0
Current Ver     8
Android Ver    3
dtype: int64
```

#Missing values in Rating should be filled with the integer value of 0

```
df_apps['Rating'] = df_apps['Rating'].fillna(int(0))
df_apps.dropna(inplace = True)
```

```
df_apps.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10829 entries, 0 to 10839
Data columns (total 13 columns):
App                10829 non-null object
Category           10829 non-null object
Rating            10829 non-null float64
Reviews           10829 non-null int64
Size               10829 non-null object
Installs           10829 non-null object
Type               10829 non-null object
Price              10829 non-null object
Content Rating     10829 non-null object
Genres             10829 non-null object
Last Updated       10829 non-null datetime64[ns]
Current Ver        10829 non-null object
Android Ver        10829 non-null object
dtypes: datetime64[ns](1), float64(1), int64(1), object(10)
memory usage: 1.2+ MB
```

Cleaning Categories into integers

```
CategoryString = df_apps["Category"]
categoryVal = df_apps["Category"].unique()
categoryValCount = len(categoryVal)
category_dict = {}
for i in range(0,categoryValCount):
    category_dict[categoryVal[i]] = i
df_apps["Category_c"] = df_apps["Category"].map(category_dict).astype(int)
```

#scaling and cleaning size of installation

```
def change_size(size):
    if 'M' in size:
        x = size[:-1]
        x = float(x)*1000000
        return(x)
    elif 'k' == size[-1:]:
        x = size[:-1]
        x = float(x)*1000
        return(x)
    else:
        return None

df_apps["Size"] = df_apps["Size"].map(change_size)
```

#filling Size which had NA

```
df_apps.Size.fillna(method = 'ffill', inplace = True)
```

#Converting Type classification into binary

```
def type_cat(types):  
    if types == 'Free':  
        return 0  
    else:  
        return 1
```

```
df_apps['Type'] = df_apps['Type'].map(type_cat)
```

#Cleaning of genres

```
GenresL = df_apps.Genres.unique()  
GenresDict = {}  
for i in range(len(GenresL)):  
    GenresDict[GenresL[i]] = i  
df_apps['Genres_c'] = df_apps['Genres'].map(GenresDict).astype(int)
```

#Cleaning prices

```
def price_clean(price):  
    if price == '0':  
        return 0  
    else:  
        price = price[1:]  
        price = float(price)  
        return price
```

```
df_apps['Price'] = df_apps['Price'].map(price_clean).astype(float)
```

convert reviews to numeric

```
df_apps['Reviews'] = df_apps['Reviews'].astype(int)
```

#Cleaning of content rating classification

```
RatingL = df_apps['Content Rating'].unique()  
RatingDict = {}  
for i in range(len(RatingL)):  
    RatingDict[RatingL[i]] = i  
df_apps['Content Rating'] = df_apps['Content Rating'].map(RatingDict).astype(int)
```

#Cleaning no of installs classification

```
df_apps['Installs'] = [int(i[:-1].replace(',','')) for i in df_apps['Installs']]
```

```
df_apps.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10829 entries, 0 to 10839
Data columns (total 15 columns):
App                10829 non-null object
Category           10829 non-null object
Rating             10829 non-null float64
Reviews            10829 non-null int32
Size               10829 non-null float64
Installs          10829 non-null int64
Type               10829 non-null int64
Price              10829 non-null float64
Content Rating     10829 non-null int32
Genres             10829 non-null object
Last Updated       10829 non-null datetime64[ns]
Current Ver        10829 non-null object
Android Ver        10829 non-null object
Category_c         10829 non-null int32
Genres_c           10829 non-null int32
dtypes: datetime64[ns](1), float64(3), int32(4), int64(2), object(5)
memory usage: 1.2+ MB

```

df_apps.describe()

	Rating	Reviews	Size	Installs	Type	Price	Content Rating	Category_c	Genres_c
count	10829.000000	1.082900e+04	1.082900e+04	1.082900e+04	10829.000000	10829.000000	10829.000000	10829.000000	10829.000000
mean	3.623197	4.446018e+05	2.186270e+07	1.547990e+07	0.073599	1.028091	0.32810	17.665343	50.466248
std	1.513263	2.929213e+06	2.252805e+07	8.507114e+07	0.261129	15.957778	0.76176	7.480582	34.489114
min	0.000000	0.000000e+00	8.500000e+03	0.000000e+00	0.000000	0.000000	0.00000	0.000000	0.000000
25%	3.700000	3.800000e+01	5.100000e+06	5.000000e+03	0.000000	0.000000	0.00000	13.000000	19.000000
50%	4.200000	2.100000e+03	1.400000e+07	1.000000e+05	0.000000	0.000000	0.00000	18.000000	38.000000
75%	4.500000	5.481500e+04	3.000000e+07	5.000000e+06	0.000000	0.000000	0.00000	23.000000	89.000000
max	5.000000	7.815831e+07	1.000000e+08	1.000000e+09	1.000000	400.000000	5.00000	32.000000	118.000000

FEATURE ENGINEERING

```
from sklearn import preprocessing
from sklearn import utils
```

Transforming the Predictive Target (Y) using Label Uncoder

```
from sklearn import preprocessing
from sklearn import utils
Y= df_apps['Rating']
X = df_apps.drop('Rating', axis=1)
lab_enc = preprocessing.LabelEncoder()
rating_encoded = lab_enc.fit_transform(Y)
print(rating_encoded)
print(utils.multiclass.type_of_target(Y))
print(utils.multiclass.type_of_target(Y.astype('int')))
print(utils.multiclass.type_of_target(rating_encoded))
```

```
[30 28 36 ...  0 34 34]
continuous
multiclass
multiclass
```

Method 1: Recursive Feature Elimination

```
# Recursive Feature Elimination
from sklearn import datasets
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestRegressor
```

```
X = df_apps.drop('Rating', axis=1)
```

```
rfc = RandomForestClassifier(random_state=101)
rfecv = RFECV(estimator=rfc, step=1, cv=StratifiedKFold(10), scoring='accuracy')
rfecv.fit(X, rating_encoded)
```

```
RFECV(cv=StratifiedKFold(n_splits=10, random_state=None, shuffle=False),
      estimator=RandomForestClassifier(bootstrap=True, class_weight=None,
                                         criterion='gini', max_depth=None,
                                         max_features='auto', max_leaf_nodes=
None,
                                         min_impurity_decrease=0.0,
                                         min_impurity_split=None,
                                         min_samples_leaf=1, min_samples_spli
t=2,
                                         min_weight_fraction_leaf=0.0,
                                         n_estimators='warn', n_jobs=None,
                                         oob_score=False, random_state=101,
                                         verbose=0, warm_start=False),
      min_features_to_select=1, n_jobs=None, scoring='accuracy', step=1,
      verbose=0)
```

```
print('Optimal number of features: {}'.format(rfecv.n_features_))
```

```
Optimal number of features: 1
```

```
X.drop(X.columns[np.where(rfecv.support_ == False)[0]], axis=1,  
inplace=True)
```

Reviews	
0	159
1	967
2	87510
3	215644
4	967

Method 2: Feature Importance

```
# Feature Importance
```

```
from sklearn import datasets
```

```
from sklearn import metrics
```

```
from sklearn.ensemble import ExtraTreesClassifier
```

```
XX = df_apps.drop('Rating', axis=1)
```

```
# fit an Extra Trees model to the data
```

```
model = ExtraTreesClassifier()
```

```
model.fit(XX, rating_encoded)
```

```
# display the relative importance of each attribute
```

```
print(model.feature_importances_)
```

```
dset = pd.DataFrame()
```

```
dset['attr'] = XX.columns
```

```
dset['importance'] = model.feature_importances_
```

```
dset = dset.sort_values(by='importance', ascending=True)
```

```
plt.figure(figsize=(9, 5))
```

```
plt.barh(y=dset['attr'], width=dset['importance'], color='#1976D2')
```


```
plt.title('RFECV - Feature Importances', fontsize=20, fontweight='bold', pad=20)
```

```
plt.xlabel('Importance', fontsize=14, labelpad=20)
```

```
plt.show()
```



```
[0.36677685 0.2721426 0.1601974 0.00272684 0.01309127 0.02086396  
0.06932111 0.09487996]
```



I will be using the Feature Importance method result that shows more important features (comparing to the recursive feature elimination option) and that I have named in the python coding “X_FI” referring to X Features Importance.

#Selected Features using Feature Importance

```
X_FI.head()
```

	Reviews	Size	Installs	Type	Price	Content Rating	Category_c	Genres_c
0	159	19000000.0	10000	0	0.0	0	0	0
1	967	14000000.0	500000	0	0.0	0	0	1
2	87510	8700000.0	5000000	0	0.0	0	0	0
3	215644	25000000.0	50000000	0	0.0	1	0	0
4	967	2800000.0	100000	0	0.0	0	0	2

DATA VISUALIZATION

App with large number of reviews

```
df_apps.loc[df_apps.Reviews == df_apps.Reviews.max()]
```

	App	Category	Rating	Reviews	Size	Installs	Type	Price	Content Rating	Genres	Last Updated	Current Ver	Android Ver	Category_c	Genres_c
2544	Facebook	SOCIAL	4.1	78158306	26000000.0	1000000000	0	0.0	1	Social	3-Aug-18	Varies with device	Varies with device	20	86

App with the largest size

```
df_apps.loc[df_apps.Size == df_apps.Size.max()]
```

	App	Category	Rating	Reviews	Size	Installs	Type	Price	Content Rating	Genres	Last Updated	Current Ver	Android Ver	Cal
	Post Bank	FINANCE	4.5	60449	100000000.0	1000000	0	0.00	0	Finance	23-Jul-18	2.9.12	4.0 and up	
	Talking Babsy Baby: Baby Games	LIFESTYLE	4.0	140995	100000000.0	10000000	0	0.00	0	Lifestyle;Pretend Play	16-Jul-18	9	4.0 and up	
	Hungry Shark Evolution	GAME	4.5	6074334	100000000.0	100000000	0	0.00	1	Arcade	25-Jul-18	6.0.0	4.1 and up	
	Mini Golf King - Multiplayer Game	GAME	4.5	531458	100000000.0	5000000	0	0.00	0	Sports	20-Jul-18	3.04.1	4.0.3 and up	
	Hungry Shark Evolution	GAME	4.5	6074627	100000000.0	100000000	0	0.00	1	Arcade	25-Jul-18	6.0.0	4.1 and up	


App with the largest num of installs

```
df_apps.loc[df_apps.Installs == df_apps.Installs.max()]
```

	App	Category	Rating	Reviews	Size	Installs	Type	Price	Content Rating	Genres	Last Updated	Current Ver	Android Ver	Category_c	Genres_c
	Google Play Books	BOOKS_AND_REFERENCE	3.9	1433233	5000000.0	1000000000	0	0.0	1	Books & Reference	3-Aug-18	Varies with device	Varies with device	3	6
	Messenger – Text and Video Chat for Free	COMMUNICATION	4.0	56642847	35000000.0	1000000000	0	0.0	0	Communication	1-Aug-18	Varies with device	Varies with device	6	10
	WhatsApp Messenger	COMMUNICATION	4.4	69119316	35000000.0	1000000000	0	0.0	0	Communication	3-Aug-18	Varies with device	Varies with device	6	10
	Google Chrome: Fast & Secure	COMMUNICATION	4.3	9642995	17000000.0	1000000000	0	0.0	0	Communication	1-Aug-18	Varies with device	Varies with device	6	10

Paid vs Free

```
sns.countplot(df_apps['Type'],label="Count")
plt.show()
```



```
df_apps.groupby('Type')['Type'].count()
```

```
Type
0    10032
1      797
Name: Type, dtype: int64
```

App which hasn't been updated

```
# App which hasn't been updated
```

```
min(df_apps['Last Updated'])
```

```
print(df_apps[df_apps['Last Updated'] == df_apps['Last Updated'].min()])
```


	App	Category	Rating	Reviews	Size	Installs
7479	FML F*ck my life + widget	FAMILY	4.2	1415	209k	100,000+

	Type	Price	Content	Rating	Genres	Last Updated	Current Ver
7479	Free	0	Everyone		Entertainment	2010-05-21	3.1

	Android Ver
7479	1.5 and up

Most popular category

```
sns.countplot(df_apps['Category'],label="Count")  
plt.xticks(rotation='vertical')  
plt.show()
```

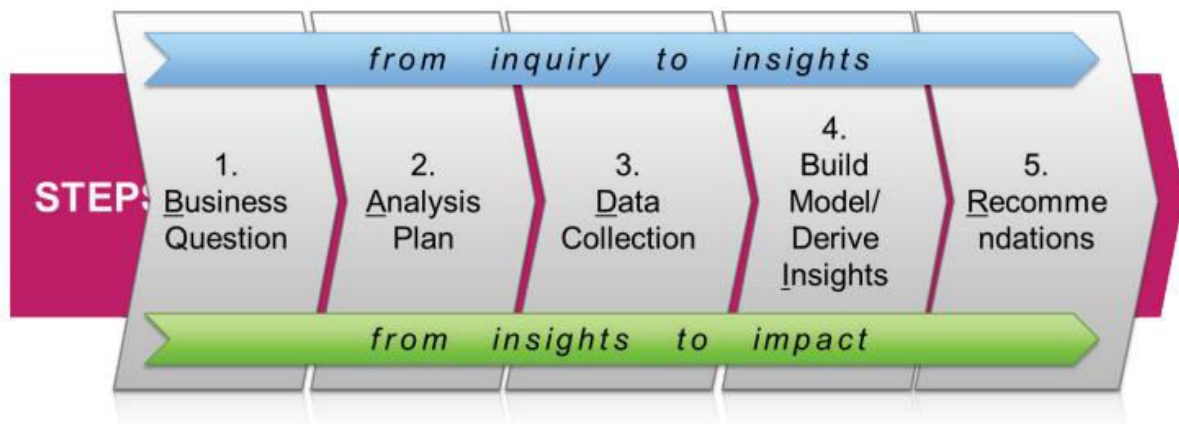


```
df_apps.groupby('Category')['Category'].count()
```

```
Category
ART_AND_DESIGN          64
AUTO_AND_VEHICLES       85
BEAUTY                   53
BOOKS_AND_REFERENCE     230
BUSINESS                 460
COMICS                   60
COMMUNICATION            387
DATING                   234
EDUCATION                156
ENTERTAINMENT            149
EVENTS                   64
FAMILY                 1968
FINANCE                  366
FOOD_AND_DRINK           127
GAME                     1144
HEALTH_AND_FITNESS       341
HOUSE_AND_HOME            88
LIBRARIES_AND_DEMO        84
LIFESTYLE                 382
MAPS_AND_NAVIGATION       137
MEDICAL                   463
NEWS_AND_MAGAZINES        283
PARENTING                 60
PERSONALIZATION           390
PHOTOGRAPHY               335
PRODUCTIVITY              424
SHOPPING                  260
SOCIAL                    295
SPORTS                    384
TOOLS                     841
TRAVEL_AND_LOCAL          258
VIDEO_PLAYERS             175
WEATHER                   82
Name: Category, dtype: int64
```

DATA SCIENCE PROCESS: BEST PRACTICES

BADIR: Project Process Framework



- Business Question

The Play Store apps data has enormous potential to drive app-making businesses to success. Actionable insights can be drawn for developers to work on and capture the Android market!

- Analysis Plan

Analysis Goal	Have a better understanding of the Android market.
Hypothesis	The social and games apps are the best rated
Methodology	Predictive Analysis Methods using Regression Algorithms like Random Forest Regression, Support Vector Regression and Linear Regression.
Project Plan	<ol style="list-style-type: none">1. Business & Data Understanding (BADIR Framework)2. Data Preparation (cleaning, Selecting and transformation)3. Modeling4. Evaluation: evaluate the result in the context of the business goal + new business requirements can pop up, due to the new patterns discovered during the data evaluation.5. Deployment: The final report needs to summarize the project insights and outcomes and review the project to see what needs to be improved upon.

- Data Collection
 - The data was scraped from the Google Play Store. Using the Web Scraping technique (for more details go back to Data Collection section in this file).
 - In order to clean and validate the data, I used the following preprocesses:
 - Remove all Missing Data
 - Feature Selection
 - Transformation of variables
 - 80% Time spent on Data Munging.
- Build Model / Derive Insights
 - Try more than one machine learning technique.
 - Fine-Tune parameters.
 - Assess Model Performance.
 - Avoid Over-fitting.
 - 20% Time spent.
- Recommendations
 - Visualizations of derived insights presenting each predictor's relationship with target "*Rating*": Histogram, scatter plots, Heat map...
 - Review existing business rules/model.
 - Target customers that are likely to default less.
 - Tracking model: Test, Measure and Improve.

PREDICTIVE MODELING AND EVALUATION (THE WHOLE PROCESS)

#XX refers to the data frame with the Important features only

```
XX.head()
depVar= df_apps['Rating']
```

#estimators

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.svm import SVR
from sklearn import linear_model
```

#model metrics

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.model_selection import cross_val_score
```

#cross validation

```
from sklearn.model_selection import train_test_split
```

#Training Set (Feature Space: X Training)

```
X_train = XX[: 1000]
X_train_count = len(X_train.index)
print('The number of observations in the Y training set are:',str(X_train_count))
X_train.head()
```

The number of observations in the Y training set are: 1000

	Reviews	Size	Installs	Type	Price	Content Rating	Category_c	Genres_c
0	159	19000000.0	10000	0	0.0	0	0	0
1	967	14000000.0	500000	0	0.0	0	0	1
2	87510	8700000.0	5000000	0	0.0	0	0	0
3	215644	25000000.0	50000000	0	0.0	1	0	0
4	967	2800000.0	100000	0	0.0	0	0	2

#Dependent Variable Training Set (y Training)

```
y_train = depVar[: 1000]
y_train_count = len(y_train.index)
#y_train_count = len(y_train)
print('The number of observations in the Y training set are:',str(y_train_count))
y_train.head()
```

```
Out[71]:
```

The number of observations in the Y training set are: 1000	
0	4.1
1	3.9
2	4.7
3	4.5
4	4.3

Name: Rating, dtype: float64

#Testing Set (X Testing)

```
X_test = XX[-100:]
X_test_count = len(X_test.index)
print('The number of observations in the feature testing set is:',str(X_test_count))
X_test.head()
```

The number of observations in the feature testing set is: 100

	Reviews	Size	Installs	Type	Price	Content Rating	Category_c	Genres_c
10740	8484	1700000.0	1000000	0	0.0	0	25	91
10741	32	7900000.0	1000	0	0.0	0	16	29
10742	16	1200000.0	500	0	0.0	0	18	19
10743	1	2000000.0	100	0	0.0	0	11	24
10744	1	5800000.0	1	0	0.0	0	11	24

#Ground Truth (y_test)

```
y_test = depVar[-100:]
y_test_count = len(y_test.index)
#y_test_count = len(y_test)
print('The number of observations in the Y training set are:',str(y_test_count))
y_test.head()
```

```
Out[73]:
```

The number of observations in the Y training set are: 100	
10740	4.2
10741	5.0
10742	3.4
10743	0.0
10744	0.0

Name: Rating, dtype: float64

```
X_train, X_test, y_train, y_test = train_test_split(X_train, y_train, test_size=0.25, random_state=0)
```

```
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
((750, 8), (750,), (250, 8), (250,))
```

#Models

```
modelSVR = SVR()
```

```
modelRF = RandomForestRegressor()
```

```
modelLR = LinearRegression()
```

Random Forest Regression

```
from sklearn.model_selection import cross_val_score
```

```
modelRF.fit(X_train,y_train)
```

```
print(cross_val_score(modelRF, X_train, y_train))
```

```
print('These values correspond to the the following: ')
```

```
print('1st value: The score array for test scores on each cv split. (Higher is an indicator of a better performing model)')
```

```
print('2nd value: The time for fitting the estimator on the train set for each cv split.')
```

```
print('3rd Value: The time for scoring the estimator on the test set for each cv split.')
```

```
print('R-Squared: %.3f' % modelRF.score(X_train,y_train))
```

```
[0.43065548 0.56608103 0.31868752]
```

```
These values correspond to the the following:
```

```
1st value: The score array for test scores on each cv split. (Higher is a  
n indicator of a better performing model
```

```
2nd value: The time for fitting the estimator on the train set for each c  
v split.
```

```
3rd Value: The time for scoring the estimator on the test set for each cv  
split.
```

```
R-Squared: 0.922
```

Support Vector Regression

```
modelSVR.fit(X_train,y_train)
```

```
print(cross_val_score(modelSVR, X_train, y_train))
```

```
print('R-Squared: %.3f' % modelSVR.score(X_train,y_train))
```

```
[-0.01695971 -0.01677 0.01178045]
```

```
R-Squared: 0.496
```

Linear Regression

```
modelLR.fit(X_train,y_train)
print(cross_val_score(modelLR, X_train, y_train))
print('R-Squared: %.3f' % modelLR.score(X_train,y_train))
```

```
[-0.01257082 -0.03493534  0.02720635]
R-Squared:  0.094
```

Predictions

#RandomForest Regression Model Predictions

```
predRF = modelRF.predict(X_test)
predRF_Rsquared = r2_score(y_test,predRF)
rmseRF = sqrt(mean_squared_error(y_test, predRF))
print('RandomForest Regression Model Predictions:')
print('RMSE: %.3f' % rmseRF)
```

#Support Vector Regression Model Predictions

```
predSVR = modelSVR.predict(X_test)
predSVR_Rsquared = r2_score(y_test,predSVR)
rmseSVR = sqrt(mean_squared_error(y_test, predSVR))
print('Support Vector Regression Model Predictions:')
print('RMSE: %.3f' % rmseSVR)
```

#Linear Regression Model Predictions

```
predLR = modelLR.predict(X_test)
predLR_Rsquared = r2_score(y_test,predLR)
rmseLR = sqrt(mean_squared_error(y_test, predLR))
print('Linear Regression Model Predictions:')
print('RMSE: %.3f' % rmseLR)
```

```
RandomForest Regression Model Predictions:
RMSE:  0.717
Support Vector Regression Model Predictions:
RMSE:  0.942
Linear Regression Model Predictions:
RMSE:  0.928
```

MODEL SELECTION

#Random Forest Selected Model with best performance

#RMSE: 0.717


#R-Squared: 0.922

```
plt.scatter(y_test, predRF, c='r', alpha = 0.5)
```

```
plt.xlabel('Ground Truth')
```

```
plt.ylabel('Predictions')
```

```
plt.show();
```



CROSS VALIDATION

Necessary imports

```
from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn import metrics
```


Perform 6-fold cross validation

```
scores = cross_val_score(modelRF, XX, depVar, cv=6)
print ("Cross-validated scores:", scores)
```

```
Cross-validated scores: [0.36185322 0.44389649 0.56824708 0.46541865 0.4687
7069 0.46052085]
```

Make cross validated predictions

```
predictions = cross_val_predict(modelRF, XX, depVar, cv=6)
plt.scatter(y_test, predRF)
```



```
accuracy = metrics.r2_score(y_test, predRF)
print ("Cross-Predicted Accuracy:", accuracy)
```

```
Cross-Predicted Accuracy: 0.45928695245513584
```